

HW4

October 29, 2019

```
In [5]: import io
import sys
import numpy as np
```

```
In [6]: vocab = np.loadtxt('./hw4_vocab.txt', dtype = str)
uni_count = np.loadtxt('./hw4_unigram.txt', dtype = float)
uni_prob = np.zeros(uni_count.shape)
total = uni_count.sum()
uni_prob = uni_count/total
vocab_m = []
m_prob = []
```

1 unigram

```
In [7]: for i in range(len(vocab)):
    if vocab[i][0] == 'M':
        #print(vocab[i])
        vocab_m.extend([vocab[i]])
        m_prob.extend([uni_prob[i]])
```

```
In [8]: for i in range(len(vocab_m)):
    print('Token: ', vocab_m[i], '      Uni-Prob: ', np.round(m_prob[i], 6))
```

```
Token: MILLION      Uni-Prob: 0.002073
Token: MORE         Uni-Prob: 0.001709
Token: MR.          Uni-Prob: 0.001442
Token: MOST         Uni-Prob: 0.000788
Token: MARKET      Uni-Prob: 0.00078
Token: MAY          Uni-Prob: 0.00073
Token: M.           Uni-Prob: 0.000703
Token: MANY         Uni-Prob: 0.000697
Token: MADE         Uni-Prob: 0.00056
Token: MUCH         Uni-Prob: 0.000515
Token: MAKE         Uni-Prob: 0.000514
Token: MONTH        Uni-Prob: 0.000445
Token: MONEY        Uni-Prob: 0.000437
```

```

Token: MONTHS      Uni-Prob: 0.000406
Token: MY           Uni-Prob: 0.0004
Token: MONDAY       Uni-Prob: 0.000382
Token: MAJOR        Uni-Prob: 0.000371
Token: MILITARY     Uni-Prob: 0.000352
Token: MEMBERS      Uni-Prob: 0.000336
Token: MIGHT        Uni-Prob: 0.000274
Token: MEETING      Uni-Prob: 0.000266
Token: MUST         Uni-Prob: 0.000267
Token: ME           Uni-Prob: 0.000264
Token: MARCH        Uni-Prob: 0.00026
Token: MAN          Uni-Prob: 0.000253
Token: MS.          Uni-Prob: 0.000239
Token: MINISTER     Uni-Prob: 0.00024
Token: MAKING        Uni-Prob: 0.000212
Token: MOVE         Uni-Prob: 0.00021
Token: MILES        Uni-Prob: 0.000206

```

2 BIGRAM

```
In [9]: bi_count = np.loadtxt('./hw4_bigram.txt',dtype = float)
```

```
In [10]: bigram_the = bi_count[bi_count[:,0]==4]
```

```
In [11]: THE_count = uni_count[3]
        THE_count
```

```
Out[11]: 3855375.0
```

```
In [14]: sorted_bigram_the[::-1][:20]
```

```
Out[14]: array([[ 4., 270.,  1.],
 [ 4., 140.,  1.],
 [ 4., 145.,  1.],
 [ 4., 139.,  1.],
 [ 4., 150.,  1.],
 [ 4., 157.,  1.],
 [ 4., 120.,  1.],
 [ 4., 164.,  1.],
 [ 4., 386.,  1.],
 [ 4.,  44.,  1.],
 [ 4., 358.,  1.],
 [ 4., 195.,  1.],
 [ 4.,  37.,  1.],
 [ 4., 200.,  1.],
 [ 4., 109.,  1.],
 [ 4., 301.,  1.]
```

```

[ 4., 343., 1.],
[ 4., 470., 1.],
[ 4., 318., 1.],
[ 4., 93., 1.]]

```

```

In [13]: sorted_bigram_the = bigram_the[bigram_the[:,2].argsort()[::-1]]
bigram_top_10 = sorted_bigram_the[:10]
top_10_vocab_index = np.subtract(bigram_top_10[:,1].astype(int),np.ones(10))
# gives us the index of the word in the 'vocab' array
top_10_words = top_10_words = vocab[top_10_vocab_index.astype(int)]
top_10_counts = bigram_top_10[:,2]
top_10_prob = top_10_counts/THE_count

```

```

In [15]: import pandas as pd
probs = {'words': top_10_words,'probability': top_10_prob}

pd.DataFrame(data = probs)

```

```

Out[15]:
   probability  words
0      0.615020  <UNK>
1      0.013372    U.
2      0.011720  FIRST
3      0.011659  COMPANY
4      0.009451    NEW
5      0.008672  UNITED
6      0.006803  GOVERNMENT
7      0.006651  NINETEEN
8      0.006287    SAME
9      0.006161    TWO

```

2.1 The stock market fell by one hundred points last week

```

In [16]: ##### UNIGRAM #####
string = 'THE STOCK MARKET FELL BY ONE HUNDRED POINTS LAST WEEK'
sentence = string.split(' ')
sentence

```

```

Out[16]: ['THE',
'STOCK',
'MARKET',
'FELL',
'BY',
'ONE',
'HUNDRED',
'POINTS',
'LAST',
'WEEK']

```

```

In [17]: total = uni_count.sum()
uni_prob = uni_count/total

```

```

vocab_ind = []
word_prob = []
word_count = []
cummu_prob = []
for i in range(len(sentence)):
    vocab_ind.extend([np.where(vocab==sentence[i])])

    prob = uni_prob[vocab_ind[i]][0]
    word_prob.extend([prob])
    word_count.extend([1])
    if i == 0:
        cummu_prob.extend([word_prob[i]])
    else:
        cummu_prob.extend([cummu_prob[i-1]*word_prob[i]])

```

```
In [18]: print('log cumulative probabiltty: ',np.log(cummu_prob[-1]))
```

```
log cumulative probabiltty: -64.50944034364878
```

```
In [257]: .047*.000668
```

```
Out[257]: 3.1395999999999996e-05
```

```
In [20]: prob_dict = {'prob': word_prob, 'word': sentence}
pd.DataFrame(columns = ['word', 'prob'], data = prob_dict)
```

```
Out[20]:
```

	word	prob
0	THE	0.047152
1	STOCK	0.000668
2	MARKET	0.000780
3	FELL	0.000265
4	BY	0.004180
5	ONE	0.006006
6	HUNDRED	0.004021
7	POINTS	0.000221
8	LAST	0.001161
9	WEEK	0.000572

```
In [21]: ##### BIGRAM APPROACH #####
string = '<s> THE STOCK MARKET FELL BY ONE HUNDRED POINTS LAST WEEK'
sentence = string.split(' ')
sentence
```

```

vocab_ind = []
bigram_num = [] #equals the vocab index +1
bigram_count = []

```

```

bigram_prob = []
word_count = []
cummu_prob = []

# first pass start logging word indexes and total occurrences

for i in range(len(sentence)):

    vocab_ind.extend([np.where(vocab==sentence[i])])
    bigram_num.extend([vocab_ind[i][0]+1])
    word_count.extend([uni_count[vocab_ind[i][0]])]

for i in range(len(sentence)-1):

    parent = bigram_num[i][0]
    child = bigram_num[i+1][0]
    cond_count = bi_count[(bi_count[:,0] == parent) & (bi_count[:,1]==child)][0][-1]
    bigram_count.extend([cond_count])
    bigram_prob.extend([cond_count/word_count[i]])
    if i == 0:
        cummu_prob.extend([bigram_prob[i]])
    else:
        cummu_prob.extend([cummu_prob[i-1]*bigram_prob[i]])

print('log likelihood:', np.log(cummu_prob[-1]))

log likelihood: [-40.91813213]

```

2.2 "The sixteen officials sold fire insurance"

UNIGRAM

```

In [22]: ##### UNIGRAM #####
string = 'THE SIXTEEN OFFICIALS SOLD FIRE INSURANCE'
sentence = string.split(' ')
sentence

total = uni_count.sum()
uni_prob = uni_count/total

vocab_ind = []
word_prob = []
word_count = []
cummu_prob = []
for i in range(len(sentence)):
    vocab_ind.extend([np.where(vocab==sentence[i])])

```

```

prob = uni_prob[vocab_ind[i]][0]
word_prob.extend([prob])
word_count.extend([])
if i == 0:
    cummu_prob.extend([word_prob[i]])
else:
    cummu_prob.extend([cummu_prob[i-1]*word_prob[i]])

print('log likelihood:', np.log(cummu_prob[-1]))

```

log likelihood: -44.291934473132606

BIGRAM

```

In [23]: string = '<s> THE SIXTEEN OFFICIALS SOLD FIRE INSURANCE'
sentence = string.split(' ')
sentence

```

```

vocab_ind = []
bigram_num = [] #equals the vocab index +1
bigram_count = []
bigram_prob = []
word_count = []
cummu_prob = []
faulty_pairs = []

```

first pass start logging word indexes and total occurrences

```

for i in range(len(sentence)):

```

```

    vocab_ind.extend([np.where(vocab==sentence[i])])
    bigram_num.extend([vocab_ind[i][0]+1])
    word_count.extend([uni_count[vocab_ind[i][0]])]

```

```

for i in range(len(sentence)-1):

```

```

    parent = bigram_num[i][0]
    child = bigram_num[i+1][0]

```

```

    try:

```

```

        cond_count = bi_count[(bi_count[:,0] == parent) & (bi_count[:,1]==child)][0][-1]
        bigram_count.extend([cond_count])
        bigram_prob.extend([cond_count/word_count[i]])

```

```

    except:

```

```

        # there is no conditional probability, child is independent of the parent,
        # calculate unigram probability for the child

        # vocab value for parent, child
        faulty_pairs.extend([(vocab[vocab_ind[i]][0], vocab[vocab_ind[i+1]][0])])
        cond_count = 0
        bigram_count.extend([cond_count])
        bigram_prob.extend([cond_count/word_count[i]])

    if i == 0:
        cummu_prob.extend([bigram_prob[i]])
    else:
        cummu_prob.extend([cummu_prob[i-1]*bigram_prob[i]])

    print('log likelihood:', np.log(cummu_prob[-1]))

log likelihood: [-inf]

```

C:\Users\Steve\Anaconda2\envs\Conda36\lib\site-packages\ipykernel_launcher.py:48: RuntimeWarning

In [26]: bigram_prob

```

Out[26]: [array([0.15865263]),
          array([0.00022851]),
          array([0.]),
          array([9.16220773e-05]),
          array([0.]),
          array([0.0030524])]

```

In []: ##### PAIRS NOTE SEEN ###

```

'SIXTEEN OFFICIALS'
'SOLD FIRE'

```

MIXTURE MODEL

```

In [31]: # for each model, the probability of each element was tallied in
        # the word_prob and bigram_prob lists
        # we will use these while tuning the lambda function
        uni_prob = word_prob # unigram model P(ith element)
        bi_prob = bigram_prob # bigram model P(child/parent)
        lmbda_score = []

        for lmbda in np.linspace(0,1.001,1000):
            prob = 1

```

```

        for i in range(len(word_prob)):


            prob = prob*(uni_prob[i]*lmbda + (1-lmbda)*bi_prob[i])

            lmbda_score.extend([prob])

In [32]: lmbda_score = np.log(lmbda_score)
        lmbda_val = np.linspace(0,1.001,1000)
        import matplotlib.pyplot as plt
        plt.xlabel('lambda')
        plt.ylabel('log likelihood')
        plt.plot(np.linspace(0,1.001,1000),lmbda_score)
        plt.show()

```

C:\Users\Steve\Anaconda2\envs\Conda36\lib\site-packages\ipykernel_launcher.py:2: RuntimeWarning:



```

In [33]: np.where(lmbda_score == np.max(lmbda_score))
        max_lambda = lmbda_val[647]
        print(max_lambda,np.max(lmbda_score))

```

0.6482952952952953 -42.9641380716605

3 4.4 Stock Market Prediction

3.1 part a

The model is essentially a time lag model that predicts the value at the next time step to be equivalent to the current time step, as evidenced by the heavy weight on a_3 . This is a poor model because it essentially admits that it does not know if the price will go up or down, so the safest bet is to predict the same as the previous step in order to minimize the prediction loss.

```
In [34]: nasdaq_train = np.loadtxt('./nasdaq00.txt', dtype = float)
        nasdaq_test  = np.loadtxt('./nasdaq01.txt', dtype = float)

        # condition the data
        nasdaq_train_three_step = []
        train_y = []
        for i in range(2, len(nasdaq_train)-1):
            nasdaq_train_three_step.append([nasdaq_train[i-2], nasdaq_train[i-1], \
                                            nasdaq_train[i]])
            train_y.extend([nasdaq_train[i+1]])

        train_y = np.array([train_y])

        nasdaq_test_three_step = []
        test_y = []
        for i in range(2, len(nasdaq_test)-1):
            nasdaq_test_three_step.append([nasdaq_test[i-2], nasdaq_test[i-1], nasdaq_test[i]])
            test_y.extend([nasdaq_test[i+1]])

        test_y = np.array([test_y])

In [36]: import numpy as np
        A_inv = np.linalg.pinv(nasdaq_train_three_step)
        weights = A_inv@train_y.T
        print(weights[:])

[[0.03189569]
 [0.01560133]
 [0.95067337]]
```

3.1.1 WEIGHTS a_1, a_2, a_3

```
In [38]: for i in range(3):
        print('a%i : %f'%(i+1, weights[i]))

a1 : 0.031896
a2 : 0.015601
a3 : 0.950673
```

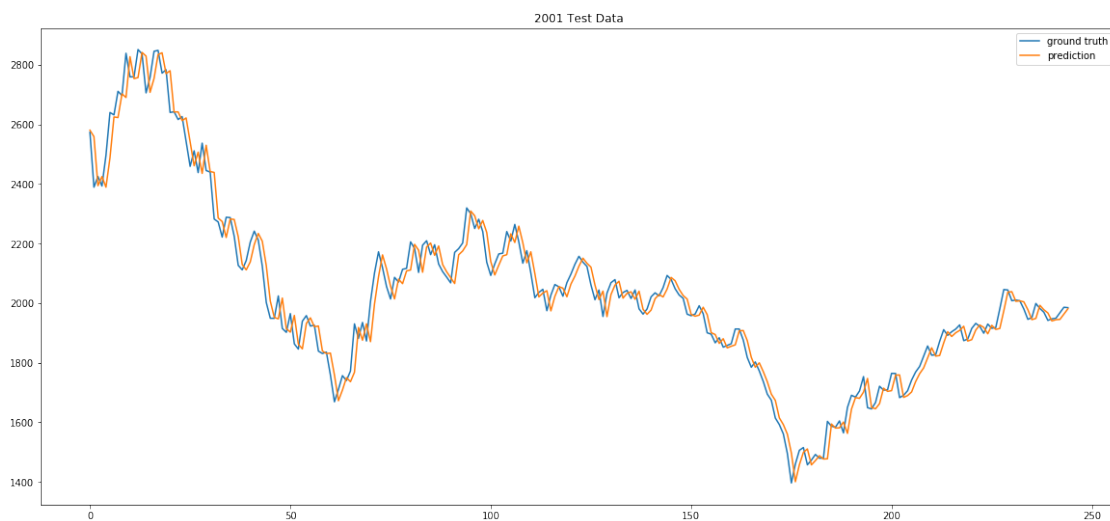

4 Part b

```
In [420]: # test the data
y_pred=[]
for i in range(test_y.shape[1]):
    pred = nasdaq_test_three_step[i]@weights
    y_pred.extend([pred])

# vector form
y_pred_tr = nasdaq_train_three_step@weights

In [419]: plt.figure(figsize=(20,20))
plt.subplot(2,1,1)
plt.plot(range(train_y.shape[1]),train_y.flatten())
plt.plot(range(train_y.shape[1]),nasdaq_train_three_step@weights)
plt.legend(['ground truth', 'prediction'])
plt.title('2001 Train Data')

plt.subplot(2,1,2)
plt.plot(range(test_y.shape[1]),test_y.flatten())
plt.plot(range(test_y.shape[1]),y_pred)
plt.legend(['ground truth', 'prediction'])
plt.title('2001 Test Data')
plt.show()
```



4.0.2 Mean Squared Error

```
In [428]: squared_error_tr = 0 # tr = train
          error_tr = 0
          squared_error_t = 0 # t = test
          error_t = 0
          for i in range(test_y.shape[1]):

              squared_error_t += (test_y[0][i]-y_pred[i])**2
              error_t += test_y[0][i]-y_pred[i]
              squared_error_tr += (train_y[0][i]-y_pred_tr[i])**2
              error_tr += train_y[0][i]-y_pred_tr[i]
```

```
mse_t = squared_error_t/test_y.shape[1]
mse_tr = squared_error_tr/train_y.shape[1]
print('MSE Test: ',np.round(mse_t[0],2),' MSE Train: ',np.round(mse_tr[0],2))
```

MSE Test: 2985.1 MSE Train: 13895.86

5 Recommendation: Don't Use

The model is essentially a time lag model that predicts the value at the next time step to be equivalent to the current time step, as evidenced by the heavy weight on a_3 . This is a poor model because it essentially admits that it does not know if the price will go up or down, so the safest bet is to predict the same as the previous step in order to minimize the prediction loss.